

Architecting for performance

A top-down approach

Ionuț Baloșin

Software Architect

www.ionutbalosin.com

 @ionutbalosin

About Me

Ionuț Baloșin

Software Architect @ Raiffeisen Bank International AG

Technical Trainer

**Java
Performance
and Tuning**


**Introduction
to Software
Architecture**

**Designing High
Performance
Applications**

www.ionutbalosin.com

 [@ionutbalosin](https://twitter.com/ionutbalosin)

Agenda



My Latency Hierarchical Model

Ultra-low Latency

(< 1ms)

Low Latency

(~ ten of ms)

Affordable Latency

(~ hundreds of ms)

Performance is not an ASR*

(~ sec)

What is **Performance**?

“Performance it’s about time and the software system’s ability to meet timing requirements.”

“Software Architecture in Practice” - Rick Kazman, Paul Clements, Len Bass



Does IT Industry Need Better Namings?



| Posted by [Ionut Balosin](#) on Apr 24, 2017. Estimated reading time: 13 minutes

Non-functional requirements or quality attributes


Oftentimes, business requirements are drilled down in functional and non-functional specifications. In this context, the non-functional requirements refer to quality attributes such as performance, availability, scalability, security, usability, etc.

But let's try to better understand what's behind the non-functional terminology guided by two approaches: the first one is to search the word in the dictionary and the second one is to show the technical implications with regards to their achievability.

According to the dictionary, the "*non*" word means "*not, absence of, unimportant, worthless*".

Based on this definition, the question which arises is: how can we name something non-functional and still having an impact on the architecture, since by semantics it is "worthless"? Why shouldn't we drop or put aside these requirements? Hence, the first inconsistency.

Following with the second approach, all non-functional requirements (e.g. security, usability) are implemented through functions (e.g. security via encryption, usability via a wizard), so they rely on concrete functions. It leads to the second question: how can we name something as being non-



DESIGN PRINCIPLES


**TACTICS, PATTERNS,
ALGORITHMS, DATA STRUCTURES**

OPERATING SYSTEM GUIDELINES

HARDWARE GUIDELINES

Cohesion


Cohesion represents the degree to which the elements inside a module work / belong together.



Classes must be cohesive, groups of class working together should be cohesive; however elements that are not related should be decoupled!

Abstractions


“The purpose of abstracting is not to be vague, but to create a new semantic level in which one can be absolutely precise” - **Edsger Dijkstra**



Abstractions => polymorphism (e.g. virtual calls) => increased runtime cost

Cyclomatic Complexity

Cyclomatic complexity is the number of linearly independent paths through a program's source code.



Higher cyclomatic complexity => branch miss predictions => pipeline stalls

Cyclomatic Complexity


Recommendation

- help the processor to make good prefetching decisions (e.g. code layout with more “predictable” branches)


But ... is it all about Big-O Complexity?

Matrix Traversal

Row traversal




Column traversal



Matrix Traversal

Row traversal



```
public long rowTraversal() {  
    long sum = 0;  
    for (int i = 0; i < mSize; i++)  
        for (int j = 0; j < mSize; j++) {  
            sum += matrix[i][j];  
        }  
  
    return sum;  
}
```


Column traversal



```
public long columnTraversal() {  
    long sum = 0;  
    for (int i = 0; i < mSize; i++)  
        for (int j = 0; j < mSize; j++) {  
            sum += matrix[j][i];  
        }  
  
    return sum;  
}
```


Matrix Traversal

Row traversal



```
public long rowTraversal() {  
    long sum = 0;  
    for (i = 0; i < mSize; i++)  
        for (j = 0; j < nSize; j++) {  
            sum += arr[i][j];  
        }  
  
    return sum;  
}
```

$O(N^2)$

Column traversal



```
public long columnTraversal() {  
    long sum = 0;  
    for (j = 0; j < nSize; j++)  
        for (i = 0; i < mSize; i++) {  
            sum += arr[i][j];  
        }  
  
    return sum;  
}
```

$O(N^2)$

Matrix Traversal

Matrix size	Row Traversal (ij) (ops/ μ s)	Column Traversal (ji) (ops/ μ s)
64 x 64	0.773	0.409
512 x 512	0.012	0.003
1024 x 1024	0.003	0.001
4096 x 4096	10^{-4}	10^{-5}
	$O(N^2)$	$O(N^2)$

higher is better

Matrix Traversal

Matrix size	Row Traversal (ij) (ops/ μ s)	Column Traversal (ji) (ops/ μ s)
64 x 64	0.773	0.409
512 x 512	0.012	0.003
1024 x 1024	0.003	0.001
4096 x 4096	10^{-4}	10^{-5}
	$O(N^2)$	$O(N^2)$

higher is better

Matrix Traversal

Matrix size	Row Traversal (ij) (ops/ μ s)	Column Traversal (ji) (ops/ μ s)
64 x 64	0.773	0.409
512 x 512	0.012	0.003
1024 x 1024	0.003	0.001
4096 x 4096	10^{-4}	10^{-5}
	$O(N^2)$	$O(N^2)$



higher is better

Why such noticeable difference ?

~ 1 order of magnitude

Matrix Traversal

Matrix size (4096 x 4096)	Row Traversal (ij)	Column Traversal (ji)
cycles per instruction	0.849	1.141
L1-dcache-loads	$10^9 \times 0.056$	$10^9 \times 9.400$
L1-dcache-load-misses	$10^9 \times 0.019$	$10^9 \times 6$
LLC-loads	$10^9 \times 0.014$	$10^9 \times 6.100$
LLC-load-misses	$10^9 \times 0.004$	$10^9 \times 0.084$
dTLB-loads	$10^9 \times 0.026$	$10^9 \times 9.400$
dTLB-load-misses	$10^3 \times 13$	$10^3 \times 101$

lower is better



Matrix Traversal

Matrix size (4096 x 4096)	Row Traversal (ij)	Column Traversal (ji)
cycles per instruction	0.849	1.141
L1-dcache-loads	$10^9 \times 0.056$	$10^9 \times 9.400$
L1-dcache-load-misses	$10^9 \times 0.019$	$10^9 \times 6$
LLC-loads	$10^9 \times 0.014$	$10^9 \times 6.100$
LLC-load-misses	$10^9 \times 0.004$	$10^9 \times 0.084$
dTLB-loads	$10^9 \times 0.026$	$10^9 \times 9.400$
dTLB-load-misses	$10^3 \times 13$	$10^3 \times 101$

lower is better

Matrix Traversal



Row traversal





NB: Simplistic representation

Matrix Traversal

Row traversal



Column traversal




NB: Simplistic representation

On modern architectures **Service Time is
highly impacted by CPU caches**

Big-O Complexity might win for huge data sets where CPU caches could not help

Recommendation

- reduce the code footprint as possible (e.g. small and clean methods)
- minimize object indirections as possible (e.g. array of primitives vs. array of objects)



DESIGN PRINCIPLES


**TACTICS, PATTERNS,
ALGORITHMS, DATA STRUCTURES**

OPERATING SYSTEM GUIDELINES

HARDWARE GUIDELINES

Caching

Caching stores application data in an optimized location to facilitate faster and easier retrieval



Data Patterns (e.g. read/write through, write behind, read ahead)


Eviction Algorithm (e.g. LRU, LFU, FIFO)

Fetching Strategy (e.g. pre-fetch, on-demand, predictive)

Topology (e.g. local, partitioned/distributed, partitioned-replicated)

Batching

Batching minimizes the number of server round trips, especially when data transfer is long.



Solution is limited by **bandwidth** and **Receiver's handling rate**

What is $\text{size}(\text{batch})$ for an optimal transfer (i.e. **max Bandwidth, min RTT**) ?

BBR Congestion Control


by Neal Cardwell, Yuchung Cheng, C. Stephen Gunn, Soheil Hassas Yeganeh, Van Jacobson

Bottleneck **B**andwidth and **R**ound-trip propagation time walk
toward (**max BW**, **min RTT**) point

[BBR Paper <https://queue.acm.org/detail.cfm?id=3022184>]

Design Asynchronous

“Design asynchronous by default, make it synchronous when it is needed” - Martin Thompson



Designing asynchronous and stateless is a good recipe for performance !

Design Asynchronous

In Java

```
java.util.concurrent.CompletableFuture<T>  
    <U> CompletableFuture<U> supplyAsync(Supplier<U> supplier)
```


```
java.util.concurrent.Future<V>  
    boolean isDone()  
    V get()
```

```
java.util.concurrent.Flow.Publisher<T>  
    void subscribe(Flow.Subscriber<? super T> subscriber)
```

```
java.util.concurrent.Flow.Subscriber<T>  
    void onSubscribe(Flow.Subscription subscription)  
    void onComplete()  
    void onNext(T item)
```

```
java.util.concurrent.Flow.Subscription  
    void request(long n)
```

Memory Access Patterns



Strided - memory access is likely to follow a predictable pattern

Spatial - nearby memory is likely to be required soon

Temporal - memory accessed recently will likely be required again soon

Memory Access Patterns


Test scenario: traverse the memory in strided, spatial and temporal fashion by accessing elements from a `long[]` array of length `2GB / sizeof(long)` (i.e. `2GB / 8`) within `4GB` of heap memory

Access Pattern	Response Time (ns / op)
Strided	0.97
Spatial	4.40
Temporal	37.34

*CPU: Intel i7-6700HQ Skylake
OS: Ubuntu 16.04.2*

Lock Free Algorithms

- ✓ failure or suspension of any thread cannot cause failure or suspension of another thread
- ✓ there is a guaranteed system-wide progress



Properties:


1. guarantees things to happen in a correct order
2. certain things happen atomically



Not very practical in the absence of hardware support (i.e. it needs state machine).

Lock Free Algorithms

Compare-And-Swap (CAS) - atomically update a memory location by another value if the previous value is the expected one



[lock] `CMPXCHG reg, reg/mem`

x86
x64



CAS APIs

j.u.c.atomic.AtomicT

boolean compareAndSet(**T** expect, **T** update)

Data Structures using CAS

j.u.c.atomic.AtomicT

T getAndIncrement()

T getAndDecrement()

T getAndAdd(**T** delta)

T getAndSet(**T** newValue)

j.u.c.locks.ReentrantLock

void lock()

boolean tryAcquire(**int** acquires)

j.u.c.ConcurrentLinkedQueue

boolean casItem(**E** cmp, **E** val)

boolean casNext(Node<**E**> cmp, Node<**E**> val)

void updateHead(Node<**E**> h, Node<**E**> p)

boolean offer(**E** e)

boolean addAll(Collection<? **extends** **E**> c)

Object Oriented Programming

```
public class Account {                                // assume sizeof(Account) ≥ 64 bytes  
    private boolean isActive;  
    private double amount;  
    private String username;  
    // ... other Fields ...  
    // ... Methods ...  
}
```

```
List<Account> allAccounts = ... // init()
```

```
for (Account account: allAccounts) { // CPU dCache inefficient layout if  
    if (!account.isActive())           // most of the accounts are active  
        triggerEvent(...);           // and JIT cannot inline it!  
}
```




dCache miss

Object Oriented Programming

```
public class Account { // assume sizeof(Account) ≥ 64 bytes
    private boolean isActive;
    private double amount;
    private String username;
    // ... other Fields ...
    // ... Methods ...
}
```

```
List<Account> allAccounts = ... // init()
```

```
for (Account account: allAccounts) { // CPU dCache inefficient layout if
    if (!account.isActive())           // most of the accounts are active
        triggerEvent(...);           // and JIT cannot inline it!
}
```



Data-Oriented Design

```
public class AccData {  
    private boolean[] areActive;  
    // ... related (e.g. used together) Fields ...  
}
```

```
AccData allAccData = ... // init()
```


```
for (int i = 0; i < accData.areActive().length; i++) {  
    if (!accData.areActive[i])  
        triggerEvent(...);  
}
```

Data-Oriented Design

```
public class AccData {  
    private boolean[] areActive;  
    // ... related (e.g. used together) Fields ...  
}
```

```
AccData allAccData = ... // init()
```


```
for (int i = 0; i < accData.areActive().length; i++) {  
    if (!accData.areActive[i])  
        triggerEvent(...);  
}
```



Data-Oriented Design

focuses on how data is read and written

- cache friendly data access patterns -



DESIGN PRINCIPLES


**TACTICS, PATTERNS,
ALGORITHMS, DATA STRUCTURES**

OPERATING SYSTEM GUIDELINES

HARDWARE GUIDELINES

Thread Affinity

Thread Affinity binds a thread to a CPU or a range of CPUs so that the thread will execute only on the designated CPU or CPUs rather than any CPU



Thread affinity takes advantages on CPU cache memory.
When a thread migrates from one processor to another all cache lines have to be moved.

Thread Affinity

In Java

<https://github.com/OpenHFT/Java-Thread-Affinity>




In Linux

taskset - retrieve or set processes's CPU affinity

sched_setaffinity(PID, ...) - sets the CPU affinity mask of the process


NUMA

Non-Uniform-Memory-Access (NUMA) is a memory design where the memory access time depends on the memory location relative to the processor




NUMA

Non-Uniform-Memory-Access (NUMA) is a memory design where the memory access time depends on the memory location relative to the processor



NUMA

Non-Uniform-Memory-Access (NUMA) is a memory design where the memory access time depends on the memory location relative to the processor



JVM NUMA-aware allocator has been implemented to take advantage of local memory

NUMA

In Java

-XX:+UseNUMA activates NUMA-aware collector

-XX:+UseParallelGC needs to be enabled as well (e.g. for Parallel Scavenger)

NUMA Performance Metrics

When evaluated against the SPEC JBB 2005 benchmark on an 8-chip Opteron machine, NUMA-aware systems showed the following performance increases:

- 32 bit – About 30 percent increase in performance with NUMA-aware allocator
- 64 bit – About 40 percent increase in performance with NUMA-aware allocator

In Linux

numactl - control NUMA policy for processes or shared memory

Workshop

[OpenJDK FAQ](#)
[Installing](#)
[Contributing](#)
[Sponsoring](#)
[Developers' Guide](#)

[Mailing lists](#)
[IRC · Wiki](#)

[Bylaws · Census](#)
[Legal](#)

JEP Process**Source code**

[Mercurial](#)
[Bundles \(6\)](#)

Groups

[\(overview\)](#)
[2D Graphics](#)
[Adoption](#)
[AWT](#)
[Build](#)
[Compatibility & Specification](#)
[Review](#)
[Compiler](#)
[Conformance](#)
[Core Libraries](#)
[Governing Board](#)
[HotSpot](#)
[Internationalization](#)
[JMX](#)

Owner Sangheon Kim

Type Feature

Scope JDK

Status Candidate

Duration M

Duplicates [JEP 157: G1 GC: NUMA-Aware Allocation](#)

Reviewed by Mikael Vidstedt, Stefan Johansson, Thomas Schatzl

Endorsed by Mikael Vidstedt

Created 2018/09/06 22:46

Updated 2018/11/10 14:26

Issue [8210473](#)

Summary


Improve G1 performance on large machines by implementing NUMA-aware memory allocation.

Non-Goals

- Support for other collectors than G1.
- Support for other OSes than Linux and Solaris.
- NUMA-awareness of other parts of the G1 collector like task queue stealing, remembered sets or refinement.

Large Pages

Using **Large Pages** the TLB can represent larger memory range hence reduces TLB misses and the number of page walks



Large Pages

In Java

-XX:+UseLargePages but it needs OS support

Operating system configuration changes to enable large pages:

- **Solaris:**

As of Solaris 9, which includes

- **Multiple Page Size Support (MPSS)**, no additional configuration is necessary.

Linux:

Large page support is included in 2.6 kernel. Some vendors have backported the code to their 2.4 based releases. To check if your system can support large page memory, try the following:

```
# cat /proc/meminfo | grep Huge
HugePages_Total: 0
HugePages_Free: 0
Hugepagesize: 2048 kB
#
```

- **Windows:**

Only Windows Server 2003 supports large page memory. In order to use it, the administrator must first assign additional privilege to the user who will be running the application:

- select Control Panel -> Administrative Tools -> Local Security Policy
- select Local Policies -> User Rights Assignment

Large Pages

Guidelines

- suitable for **intensive memory** applications with **large contiguous memory accesses**



Enable Large Pages when number of TLB misses and TLB Page walk take a significant amount of time (i.e. `dtlb_load_misses_*` CPU counters)

Large Pages

Not Recommended for ...

- short lived applications with **small working set**
- applications with **large** but **sparsely used heap**

RamFS & TmpFS

RamFS & TmpFS allocate a part of the physical memory to be used as a partition (e.g. write/read files).



Useful for applications which performs a lot disk reads/writes (e.g. logging, auditing)


RamFS & TmpFS

Test scenario: sequentially reading/writing 8GB in chunks of 4KB/512 KB on HDD/SSD/RAMFS

Chunk	HDD (5,400RPM)		SSD		RAMFS	
	Read MB/s	Write MB/s	Read MB/s	Write MB/s	Read MB/s	Write MB/s
4K	128	99	964	742	7,971	4,420
512K	147	113	1,021	788	10,760	6,045

higher is better

NB: higher read rates are caused by buffers/caches effect



DESIGN PRINCIPLES

**TACTICS, PATTERNS,
ALGORITHMS, DATA STRUCTURES**

OPERATING SYSTEM GUIDELINES

HARDWARE GUIDELINES

False Sharing

False Sharing is purely a CPU Cache issue

```
public class FalseSharing {  
    public int X;  
    public int Y;  
}
```

```
FalseSharing sharedInstance = new FalseSharing();
```

Thread 1




```
void incrementX () {  
    sharedInstance.X ++  
}
```

Thread 2



```
void incrementY () {  
    sharedInstance.Y ++  
}
```



-----➔ Request for Ownership (I -> M)
-----➔ Update

False Sharing

In Java

@Contended annotation pads fields so they will sit on different cache lines

```
public class FalseSharing {  
    public int X;  
  
    @sun.misc.Contended  
    public int Y;  
}
```

Threads	Baseline (# of ops / ms)	@Contended (# of ops / ms)
2	478.051	1,378.564

higher is better

False Sharing


Guidelines

- independent values sits on the **same cache line**
- different cores **concurrently access** that line
- there is at **least one writer** thread
- **high frequency** of writing/reading

Solid State Drive


TRIM
ON | OFF

I/O Scheduler
NOOP | Deadline | CFQ




Solid State Drive

Test scenario: sequentially writing/reading 32GB in chunks of 512 KB on SSD



Solid State Drive

Test scenario: sequentially writing/reading 32GB in chunks of 512 KB on SSD



My Latency Hierarchical Model

Ultra-low Latency

Thread affinity, NUMA, large pages, false sharing, Data-Oriented Design, CPU caches

Low Latency

Memory access patterns, lock free, asynchronous processing, stateless, RamFS/TmpFS

Affordable Latency

Algorithms complexities, data structures, batching, caching

Performance is not an ASR

Small and clean methods, cyclomatic complexity, cohesion, abstractions

NB: Model is not exclusive and might be subject of changes

*Performance is simple, you just have to be
aware of everything!*

Ionuț Baloșin



THANK YOU

Ionuț Baloșin


Software Architect

www.ionutbalosin.com



@ionutbalosin

Further References



Articles by **Ulrich Drepper**

- ✓ What every programmer should know about memory
- ✓ CPU caches
- ✓ Virtual memory
- ✓ NUMA systems
- ✓ What programmers can do - cache optimization
- ✓ What programmers can do - multi-threaded optimizations
- ✓ Memory performance tools

Further References

- ✓ **Performance Methodology Mindmap - Kirk Perpendine and Alexey Shipilev**
 - <https://shipilev.net/talks/devoxx-Nov2012-perfMethodology-mindmap.pdf>
- ✓ **Cpu Caches and Why You Care - Scott Meyers**
- ✓ **CPU caches - Ulrich Drepper**
- ✓ **Async or Bust!? - Todd Montgomery**
- ✓ <http://mechanical-sympathy.blogspot>
- ✓ **An Introduction to Lock-Free Programming**
 - <http://preshing.com/20120612/an-introduction-to-lock-free-programming>
- ✓ **Intel's 'cmpxchg' instruction**
 - <http://heather.cs.ucdavis.edu/~matloff/50/PLN/lock.pdf>
- ✓ <http://docs.oracle.com/javase/7/docs/technotes/guides/vm/performance-enhancements-7.html>
- ✓ <http://www.thegeekstuff.com/2008/11/overview-of-ramfs-and-tmpfs-on-linux>

THANK YOU

Ionuț Baloșin

Software Architect

www.ionutbalosin.com



@ionutbalosin