



Runtime vs. compile time (JIT vs AOT) optimizations in Java and C++

Ionuț Baloșin

Software Architect

www.ionutbalosin.com



@ionutbalosin



Krakow, 20-22 June 2018

Copyright © 2018 by Ionuț Baloșin

@ionutbalosin

About Me

Ionuț Baloșin

Software Architect @ LUXOFT

Technical Trainer

- Java Performance and Tuning
- Introduction to Software Architecture
- Designing High Performance Applications



www.ionutbalosin.com

 @ionutbalosin

Agenda

Sequential sum of N elements array ($\sum_{i=1}^N \text{array}[i]$)

Sequential sum of N integers ($\sum_{i=1}^N i$)

Fields Layout

Null Checks


Lock Elision

Virtual Calls

Scalar Replacement

Conclusions

Talk Motivation



Talk Motivation

The “Pleasure of Finding Things Out”

Learning based on evidence driven by
experiments



Key Points

Reveal optimizations generated at runtime by [JIT C2](#) in comparison to compile time optimizations triggered by [LLVM clang](#).

How powerful* are runtime (i.e. [JIT C2](#)) vs. ahead of time (i.e. [LLVM clang](#)) optimizations?

This presentation is **not a battle**, I **do not** try to **establish a winner**, but just to **study different optimization approaches**!

powerful - not necessary linked to timings !*

Why it might be a matter of interest?

Performance

LLVM Clang is one of the best AOT Compilers

JIT C2 is still the most popular option for Java

Java is moving towards AOT (JEP-295)

Compilation Process

JIT

Java Source Code

Java Bytecode

JVM

Interpreter

Profiler

JIT Compiler

Native Code

JDK 10 / HotSpot C2

clang

C++ Source Code

Clang Front End

LLVM Optimizer

Code Generator

Native Code

C++17 / clang 6 w/ -O3

Methodology and Pitfalls

Comparing LLVM clang vs. JIT C2 optimizations using different programming languages is difficult:

- it is not always 🍏 to 🍏 (e.g. not the same source code)
- language specifics (e.g. Java, C++)
- benchmarks specifics (e.g. Java Microbenchmark Harness, C++ Google Benchmark)
- run on single machine and report (e.g. i7-6700HQ Skylake w/ 16GB DDR4)
- run using single Java compiler (e.g. JIT C2) and report
- “fast enough is fast enough” - Cliff Click

Sequential sum of N elements array

$$(\sum_{i=1}^N array[i])$$

Sequential sum of N elements array

Source code

```
private int[] array;

@Benchmark
public long hotMethod() {
    long sum = 0;

    for (int i = 0; i < array.length; i++) {
        sum += array[i];
    }

    return sum;
}
```

Sequential sum of N elements array

JIT C2

```
mov r11d,DWORD PTR [rsi+0xc]      ;*getfield array
mov r10d,DWORD PTR [r12+r11*8+0xc] ;*arraylength

mov eax,DWORD PTR [r12+r11*8+0x10] ;*iaload
nop DWORD PTR [rax+rax*1+0x0]      ;*iload_1
```

```
+ L0:    add eax,DWORD PTR [r8+rcx*4+0x10]
|        add eax,DWORD PTR [r8+rcx*4+0x14]
|        add eax,DWORD PTR [r8+rcx*4+0x18]
|        add eax,DWORD PTR [r8+rcx*4+0x1c]
|        ...
|        add eax,DWORD PTR [r8+rcx*4+0x40]
|        add eax,DWORD PTR [r8+rcx*4+0x44]
|        add eax,DWORD PTR [r8+rcx*4+0x48]
|        add eax,DWORD PTR [r8+rcx*4+0x4c] ;*iadd
|        add ecx,0x10                      ;*iinc
|        cmp ecx,r9d
+        jl L0                            ;*if_icmpge

+ L1:    add eax,DWORD PTR [r8+rcx*4+0x10] ;*iadd
|        inc ecx                          ;*iinc
|        cmp ecx,r10d
+        jl L1                            ;*if_icmpge
```

Sequential sum of N elements array

JIT C2

scalar pre loop
- 1st integer add -

```
mov r11d,DWORD PTR [rsi+0xc]      ;*getfield array
mov r10d,DWORD PTR [r12+r11*8+0xc] ;*arraylength

mov eax,DWORD PTR [r12+r11*8+0x10] ;*iaload
nop DWORD PTR [rax+rax*1+0x0]      ;*iload_1

+ L0: add eax,DWORD PTR [r8+rcx*4+0x10]
      | add eax,DWORD PTR [r8+rcx*4+0x14]
      | add eax,DWORD PTR [r8+rcx*4+0x18]
      | add eax,DWORD PTR [r8+rcx*4+0x1c]
      | ...
      | add eax,DWORD PTR [r8+rcx*4+0x40]
      | add eax,DWORD PTR [r8+rcx*4+0x44]
      | add eax,DWORD PTR [r8+rcx*4+0x48]
      | add eax,DWORD PTR [r8+rcx*4+0x4c] ;*iadd
      | add ecx,0x10 ;*iinc
      | cmp ecx,r9d
+     jl L0 ;*if_icmpge

+ L1: add eax,DWORD PTR [r8+rcx*4+0x10] ;*iadd
      | inc ecx ;*iinc
      | cmp ecx,r10d
+     jl L1 ;*if_icmpge
```

Sequential sum of N elements array

JIT C2

scalar pre loop
- 1st integer add -

```
mov r11d,DWORD PTR [rsi+0xc]      ;*getfield array
mov r10d,DWORD PTR [r12+r11*8+0xc] ;*arraylength

mov eax,DWORD PTR [r12+r11*8+0x10] ;*iaload
nop DWORD PTR [rax+rax*1+0x0]      ;*iload_1
```

scalar main loop
loop unrolling
- 16 integer adds per cycle -

```
+ L0:  add eax,DWORD PTR [r8+rcx*4+0x10]
      |  add eax,DWORD PTR [r8+rcx*4+0x14]
      |  add eax,DWORD PTR [r8+rcx*4+0x18]
      |  add eax,DWORD PTR [r8+rcx*4+0x1c]
      |  ...
      |  add eax,DWORD PTR [r8+rcx*4+0x40]
      |  add eax,DWORD PTR [r8+rcx*4+0x44]
      |  add eax,DWORD PTR [r8+rcx*4+0x48]
      |  add eax,DWORD PTR [r8+rcx*4+0x4c] ;*iadd
      |  add ecx,0x10                      ;*iinc
      |  cmp ecx,r9d
      |  +
      |  jl L0                            ;*if_icmpge

+ L1:  add eax,DWORD PTR [r8+rcx*4+0x10] ;*iadd
      |  inc ecx                        ;*iinc
      |  cmp ecx,r10d
      |  +
      |  jl L1                            ;*if_icmpge
```

Sequential sum of N elements array

JIT C2

scalar pre loop
- 1st integer add -

```
mov r11d,DWORD PTR [rsi+0xc] ;*getfield array
mov r10d,DWORD PTR [r12+r11*8+0xc] ;*arraylength
```

```
mov eax,DWORD PTR [r12+r11*8+0x10] ;*iaload
nop DWORD PTR [rax+rax*1+0x0] ;*iload_1
```

scalar main loop
loop unrolling
- 16 integer adds per cycle -

```
+ L0:
| add eax,DWORD PTR [r8+rcx*4+0x10]
| add eax,DWORD PTR [r8+rcx*4+0x14]
| add eax,DWORD PTR [r8+rcx*4+0x18]
| add eax,DWORD PTR [r8+rcx*4+0x1c]
| ...
| add eax,DWORD PTR [r8+rcx*4+0x40]
| add eax,DWORD PTR [r8+rcx*4+0x44]
| add eax,DWORD PTR [r8+rcx*4+0x48]
| add eax,DWORD PTR [r8+rcx*4+0x4c] ;*iadd
| add ecx,0x10 ;*iinc
| cmp ecx,r9d
|
+ jnl L0 ;*if_icmpge
```

scalar post loop
- 1 integer add per cycle -

```
+ L1:
| add eax,DWORD PTR [r8+rcx*4+0x10] ;*iadd
| inc ecx ;*iinc
| cmp ecx,r10d
|
+ jnl L1 ;*if_icmpge
```

Sequential sum of N elements array

LLVM clang

```
+ L0:  movdqu xmm2, xmmword ptr [rcx + 4*rdi]
      |      paddq xmm2, xmm0
      |      movdqu xmm0, xmmword ptr [rcx + 4*rdi + 16]
      |      paddq xmm0, xmm1
      |      ...
      |      movdqu xmm0, xmmword ptr [rcx + 4*rdi + 96]
      |      paddq xmm0, xmm4
      |      movdqu xmm1, xmmword ptr [rcx + 4*rdi + 112]
      |      paddq xmm1, xmm2
      |      add rdi, 32
      |      add rdx, 4
+      jne L0

+ L1:  movdqu xmm2, xmmword ptr [rdx - 16]
      |      paddq xmm0, xmm2
      |      movdqu xmm2, xmmword ptr [rdx]
      |      paddq xmm1, xmm2
      |      add rdx, 32
      |      inc rax
+      jne L1

+ L2:  add eax, dword ptr [rcx + 4*rsi]
      |      inc rsi
      |      cmp rsi, r9
+      jb L2
```


Sequential sum of N elements array

LLVM clang

vectorized main loop
loop unrolling
- 8x4 adds per cycle -

```
+ L0:  movdqu xmm2, xmmword ptr [rcx + 4*rdi]
      |  padd  xmm2, xmm0
      |  movdqu xmm0, xmmword ptr [rcx + 4*rdi + 16]
      |  padd  xmm0, xmm1
      |  ...
      |  movdqu xmm0, xmmword ptr [rcx + 4*rdi + 96]
      |  padd  xmm0, xmm4
      |  movdqu xmm1, xmmword ptr [rcx + 4*rdi + 112]
      |  padd  xmm1, xmm2
      |  add   rdi, 32
      |  add   rdx, 4
      |  jne   L0
+
+ L1:  movdqu xmm2, xmmword ptr [rdx - 16]
      |  padd  xmm0, xmm2
      |  movdqu xmm2, xmmword ptr [rdx]
      |  padd  xmm1, xmm2
      |  add   rdx, 32
      |  inc   rax
      |  jne   L1
+
+ L2:  add   eax, dword ptr [rcx + 4*rsi]
      |  inc   rsi
      |  cmp   rsi, r9
      |  jbe   L2
+
```

[note] XMM - 128 bit register
[note] 1 add - 1 array integer add

Sequential sum of N elements array

LLVM clang

vectorized main loop
loop unrolling
- 8x4 adds per cycle -

```
+ L0:  movdqu xmm2, xmmword ptr [rcx + 4*rdi]
      |  paddq xmm2, xmm0
      |  movdqu xmm0, xmmword ptr [rcx + 4*rdi + 16]
      |  paddq xmm0, xmm1
      |  ...
      |  movdqu xmm0, xmmword ptr [rcx + 4*rdi + 96]
      |  paddq xmm0, xmm4
      |  movdqu xmm1, xmmword ptr [rcx + 4*rdi + 112]
      |  paddq xmm1, xmm2
      |  add rdi, 32
      |  add rdx, 4
      |  jne L0
```

vectorized post loop
loop unrolling
- 2x4 adds per cycle -

```
+ L1:  movdqu xmm2, xmmword ptr [rdx - 16]
      |  paddq xmm0, xmm2
      |  movdqu xmm2, xmmword ptr [rdx]
      |  paddq xmm1, xmm2
      |  add rdx, 32
      |  inc rax
      |  jne L1
```

```
+ L2:  add eax, dword ptr [rcx + 4*rsi]
      |  inc rsi
      |  cmp rsi, r9
      |  jb L2
```

[note] XMM - 128 bit register
[note] 1 add - 1 array integer add

Sequential sum of N elements array

LLVM clang

vectorized main loop
loop unrolling
- 8x4 adds per cycle -

```
+ L0:  movdqu xmm2, xmmword ptr [rcx + 4*rdi]
      |  paddq xmm2, xmm0
      |  movdqu xmm0, xmmword ptr [rcx + 4*rdi + 16]
      |  paddq xmm0, xmm1
      |  ...
      |  movdqu xmm0, xmmword ptr [rcx + 4*rdi + 96]
      |  paddq xmm0, xmm4
      |  movdqu xmm1, xmmword ptr [rcx + 4*rdi + 112]
      |  paddq xmm1, xmm2
      |  add rdi, 32
      |  add rdx, 4
      |  jne L0
```

vectorized post loop
loop unrolling
- 2x4 adds per cycle -

```
+ L1:  movdqu xmm2, xmmword ptr [rdx - 16]
      |  paddq xmm0, xmm2
      |  movdqu xmm2, xmmword ptr [rdx]
      |  paddq xmm1, xmm2
      |  add rdx, 32
      |  inc rax
      |  jne L1
```

scalar post loop
- 1 add per cycle -

```
+ L2:  add eax, dword ptr [rcx + 4*rsi]
      |  inc rsi
      |  cmp rsi, r9
      |  jb L2
```

[note] XMM - 128 bit register
[note] 1 add - 1 array integer add

Sequential sum of N elements array

JIT C2

scalar pre loop

scalar main loop

scalar post loop



16 integers

[note] not general purpose pattern !

□ 32 bit

□ □ □ □ 128 bit

Sequential sum of N elements array

JIT C2

scalar pre loop

scalar main loop

scalar post loop



16 integers

[note] not general purpose pattern !

LLVM clang

vectorized main loop

vectorized post loop

scalar post loop



32 integers

8 integers

□ 32 bit

□□□□ 128 bit

Sequential sum of N elements array

Benchmark comparison

Array Size	JIT C2 us/op	LLVM Clang us/op
10×10^3	3.2	0.8
100×10^3	32.8	9.9
$1\,000 \times 10^3$	943.3	297.7

lower is better

[note] 1 op – 1 method call

Sequential sum of N elements array

Conclusions

LLVM clang uses vectorization for:

- **main loop** - 5 XMM SSE registers on 128 bits → 32 integer adds per loop
- **post loop** - 3 XMM SSE registers on 128 bits → 8 integer adds per loop

LLVM clang does the **remaining post loop** without unrolling and without vectorization, just adding one by one

Optimizations done by JIT C2 (e.g. loop peeling, loop unrolling) without vectorization support could not outperform LLVM clang



JDK / JDK-8188313

C2: Consider enabling auto-vectorization for simple reductions

Details

Type:	 Enhancement	Status:	OPEN
Priority:	3 P3	Resolution:	Unresolved
Affects Version/s:	9	Fix Version/s:	tbd_feature
Component/s:	hotspot		
Labels:	c2 c2-loopopts community-candidate performance		
Subcomponent:	compiler		

Description

Reconsider the decision to disable autovectorization for simple reductions made by [JDK-8078563](#).

Richard Startin reports [1] that artificially complicating reduction operation (to make it eligible for auto-vectorization) produces a better result than original code (sum of elements).

**But ... which are the cases when JIT does
loop vectorization?**

Sum(array[], array[])

```
private int[] a, int[] b, int[] c;
private int len;

@Setup(Level.Trial)
public void setUp() {
    this.len = 100_000_000;
    this.a = new int[len];
    this.b = new int[len];
    this.c = new int[len];
    for (int i = 0; i < len; i++) {
        a[i] = i + 1;
        b[i] = i + 1;
    }
}

@Benchmark
public int[] hotMethod() {
    for (int i = 0; i < len; i++) {
        c[i] = a[i] + b[i];
    }
    return c;
}
```

Sum(array[], array[])

JIT C2

vectorized main loop
loop unrolling
- 16x8 adds per cycle -

```
+ L0:  vmovdqu ymm0, YMMWORD PTR [r9+rbx*4+0x10]
|      vpaddq ymm0, ymm0, YMMWORD PTR [rax+rbx*4+0x10]
|      vmovdqu YMMWORD PTR [rbp+rbx*4+0x10], ymm0
|      movsxd r10, ebx
|      vmovdqu ymm0, YMMWORD PTR [r9+r10*4+0x30]
|      vpaddq ymm0, ymm0, YMMWORD PTR [rax+r10*4+0x30]
|      vmovdqu YMMWORD PTR [rbp+r10*4+0x30], ymm0
|      vmovdqu ymm0, YMMWORD PTR [r9+r10*4+0x50]
|      vpaddq ymm0, ymm0, YMMWORD PTR [rax+r10*4+0x50]
|      vmovdqu YMMWORD PTR [rbp+r10*4+0x50], ymm0
|      ...
|      vmovdqu ymm0, YMMWORD PTR [r9+r10*4+0x1d0]
|      vpaddq ymm0, ymm0, YMMWORD PTR [rax+r10*4+0x1d0]
|      vmovdqu YMMWORD PTR [rbp+r10*4+0x1d0], ymm0
|      vmovdqu ymm0, YMMWORD PTR [r9+r10*4+0x1f0]
|      vpaddq ymm0, ymm0, YMMWORD PTR [rax+r10*4+0x1f0]
|      vmovdqu YMMWORD PTR [rbp+r10*4+0x1f0], ymm0
|      add ebx, 0x80
|      cmp ebx, esi
+      jl L0
```

[note] YMM - 256 bit register
[note] 1 add - 1 array integer add

Sum(array[], const)

```
@Param({"5"})
public Integer anInt;

private int[] array;

@Setup(Level.Trial)
public void setUp() {
    this.array = new int[100_000_000];
    for (int i = 0; i < array.length; i++)
        array[i] = i + 1;
}

@Benchmark
public int [] hotMethod() {
    for (int i = 0; i < array.length; i++) {
        array[i] = array[i] + anInt.intValue();
    }
    return array;
}
```

Sum(array[], const)

JIT C2

vectorized main loop
loop unrolling
- 16x8 adds per cycle -

```
+ L0:
| vpaddq ymm0,ymm1,YMMWORD PTR [rdx+rsi*4+0x10]
| vmovdqu YMMWORD PTR [rdx+rsi*4+0x10],ymm0
| vpaddq ymm0,ymm1,YMMWORD PTR [rdx+rsi*4+0x30]
| vmovdqu YMMWORD PTR [rdx+rsi*4+0x30],ymm0
| vpaddq ymm0,ymm1,YMMWORD PTR [rdx+rsi*4+0x50]
| vmovdqu YMMWORD PTR [rdx+rsi*4+0x50],ymm0
|
| ...
| vpaddq ymm0,ymm1,YMMWORD PTR [rdx+rsi*4+0x1b0]
| vmovdqu YMMWORD PTR [rdx+rsi*4+0x1b0],ymm0
| vpaddq ymm0,ymm1,YMMWORD PTR [rdx+rsi*4+0x1d0]
| vmovdqu YMMWORD PTR [rdx+rsi*4+0x1d0],ymm0
| vpaddq ymm0,ymm1,YMMWORD PTR [rdx+rsi*4+0x1f0]
| vmovdqu YMMWORD PTR [rdx+rsi*4+0x1f0],ymm0
| add esi,0x80
| cmp esi,r8d
+ jl L0
```

[note] YMM - 256 bit register
[note] 1 add - 1 array integer add

JIT Vectorization Loop


scalar pre loop

e.g. memory alignment, CPU caches



vectorized main loop

e.g. 16x8 integer adds per cycle



vectorized post loop

e.g. 8 integer adds per cycle



scalar post loop

e.g. 1 integer add per cycle



[note] pattern valid for JDK 10

Benchmark comparison

Test	Array Size	JIT C2 ms/op	LLVM Clang ms/op
Sum(array[], array[])	10 ⁹	61.9	62.4
Sum(array[], const)	10 ⁹	30.0	31.6

lower is better

“Fast enough is fast enough” - Cliff Click

Sequential sum of N integers

$$(\sum_{i=1}^N i)$$

Sequential sum (N integers)

Source code

```
@Benchmark
public long hotMethod() {
    long sum = 0;

    for (int i = 1; i <= N; i++) {
        sum += i;
    }

    return sum;
}
```

Sequential sum (N integers)

JIT C2

```
mov r10d,ebx
inc r10d

+L0:  movsxd r11,ebx
      add rbp,r11
      inc ebx
      cmp ebx,r10d
+    jl L0
      jmp L2

+L1:  mov rbp,rax
      L2: movsxd r10,ebx
      add rbp,r10
      ...
      mov rax,r10
      add rax,rbp
      add rax,r10
      ...
      add rax,0x78
      add ebx,0x10
      cmp ebx,0x5f5e0f2
+    jl L1

+L3:  movsxd r10,ebx
      add rax,r10
      inc ebx
      cmp ebx,0x5f5e101
+    jl L3
```

Sequential sum (N integers)

JIT C2

scalar pre loop
- 1 integer add per cycle -

+L0:

$$+ \quad |$$

```
mov  r10d, ebx
inc  r10d
```

```
movsxd r11,ebx
add rbp,r11
inc ebx
cmp ebx,r10d
jl L0
jmp L2
```

+L1:

+

```

rbp, rax
L2: movsxd r10, ebx
add rbp, r10
...
mov rax, r10
add rax, rbp
add rax, r10
...
add rax, 0x78
add ebx, 0x10
cmp ebx, 0x5f5e0f2
jl L1

```

+L3:

+

```
movsxd r10,ebx
add rax,r10
inc ebx
cmp ebx,0x5f5e101
jl L3
```

Sequential sum (N integers)

JIT C2

scalar pre loop
- 1 integer add per cycle -

+L0:

+

```
mov r10d,ebx
inc r10d
```

```
movsxd r11,ebx
add rbp,r11
inc ebx
cmp ebx,r10d
jl L0
jmp L2
```

scalar main loop
loop unrolling
- 16 integer adds per cycle -

+L1:

+

```
mov rbp,rax
L2: movsxd r10,ebx
add rbp,r10
...
mov rax,r10
add rax,rbp
add rax,r10
...
add rax,0x78
add ebx,0x10
cmp ebx,0x5f5e0f2
jl L1
```

+L3:

+

```
movsxd r10,ebx
add rax,r10
inc ebx
cmp ebx,0x5f5e101
jl L3
```

Sequential sum (N integers)

JIT C2

scalar pre loop
- 1 integer add per cycle -

+L0:

+

```
mov r10d,ebx
inc r10d
```

```
movsxd r11,ebx
add rbp,r11
inc ebx
cmp ebx,r10d
jl L0
jmp L2
```

scalar main loop
loop unrolling
- 16 integer adds per cycle -

+L1:

+

```
mov rbp,rax
L2: movsxd r10,ebx
add rbp,r10
...
mov rax,r10
add rax,rbp
add rax,r10
...
add rax,0x78
add ebx,0x10
cmp ebx,0x5f5e0f2
jl L1
```

scalar post loop
- 1 integer add per cycle -

+L3:

+

```
movsxd r10,ebx
add rax,r10
inc ebx
cmp ebx,0x5f5e101
jl L3
```

Sequential sum (N integers)

LLVM clang

```
mov    eax, edi
dec    edi
imul   rdi, rax
shr    rdi
add    rdi, rax
mov    rax, rdi
```

Reduction formula: $\frac{N * (N - 1)}{2} + N$

Sequential sum (N integers)

LLVM clang

no loop

- induction variable optimization -

```
mov    eax, edi
dec    edi
imul   rdi, rax
shr    rdi
add    rdi, rax
mov    rax, rdi
```

Reduction formula: $\frac{N * (N - 1)}{2} + N$

Sequential sum (N integers)

Benchmark comparison

N size	JIT C2 ms/op	LLVM Clang ms/op
10^6	0.3	2×10^{-6}
10×10^6	3.1	2×10^{-6}
100×10^6	31.1	2×10^{-6}

lower is better

Sequential sum (N integers)

Conclusions

LLVM clang is able to recognize this pattern and to optimize it to a *reduction formula*, hence the throughput is almost constant

Runtime performance is not dependent on size of **N**.

Optimizations done by JIT C2 (e.g. loop peeling, loop unrolling) do not bring too much of improvements in comparison to LLVM clang.

Runtime performance is dependent on size of **N**.

Sequential sum (N integers)


Hypothetically, **JIT** could also optimize the whole loop into a reduction formula (e.g. $\frac{N*(N-1)}{2}$), but **it would need more advanced loop analysis** (e.g. something akin to LLVM's **Scalar Evolution and Loop Optimization** - to recognize how sum is related to the *loop induction variable*) **which at the moment might not happen in case of JIT.**

Fields Layout

Fields Layout


JAVA

```
class Vanilla {  
    short x;  
    short y;  
}
```



C++


```
struct Vanilla {  
    short x;  
    short y;  
}
```



 `h` object header

[note] 64-bit object header / compressed OOPs / 8 bytes aligned

Fields Layout



	# of objects fits Cache Line (Java class – 16B)	# of objects fits Cache Line (C++ struct – 4B)
Vanilla	4	16

Fields Layout

```
@Benchmark
@Group("Vanilla")
public double testVanilla(VanillaState state) {
    long sum = 0;
    for (int i = 0; i < state.array.length; i++) {
        sum += (state.array[i]).x + (state.array[i]).y; // vanilla.x + vanilla.y
    }
    return sum;
}
```

```
/** array.length = 2048 */
```

Fields Layout

Benchmark comparison


	JIT C2 us/op	LLVM Clang us/op
Vanilla	1.32	1.18

lower is better

Fields Layout


JAVA

```
class Padding {  
    short x;  
    long 10, 11, 12, 13, 14, 15, 16, 17;  
    short y;  
}
```



C++

```
struct Padding {  
    short x;  
    short y;  
    long 10, 11, 12, 13, 14, 15, 16, 17;  
}
```




object header



alignment/padding

[note] 64-bit object header / compressed OOPs / 8 bytes aligned

Fields Layout



	# of objects fits Cache Line (Java class – 80B)	# of objects fits Cache Line (C++ struct – 72B)
Padding	1 (X and Y sit on the same cache line)	1 (X and Y sit on the same cache line)

Fields Layout

```
@Benchmark
@Group("Padding")
public double testPadding(PaddingState state) {
    long sum = 0;
    for (int i = 0; i < state.array.length; i++) {
        sum += (state.array[i]).x + (state.array[i]).y; // padding.x + padding.y
    }
    return sum;
}
```

```
/** array.length = 2048 */
```

Fields Layout

Benchmark comparison


	JIT C2 us/op	LLVM Clang us/op
Padding	1.74	1.73

lower is better

Fields Layout


JAVA

```
class Base {  
    short x;  
    long 10, 11, 12, 13, 14, 15, 16, 17;  
}  
class Hierarchy extends Base {  
    short y;  
}
```



C++

```
struct Base {  
    short x;  
    long 10, 11, 12, 13, 14, 15, 16, 17;  
}  
struct Hierarchy: Base {  
    short y;  
}
```




object header



alignment/padding

[note] 64-bit object header / compressed OOPs / 8 bytes aligned

Fields Layout



of objects fits Cache Line
(Java class – 88B)

Hierarchy

1 (**X** and **Y** sit on different cache lines)

of objects fits Cache Line
(C++ struct – 80B)

1 (**X** and **Y** sit on different cache lines)

Fields Layout

```
@Benchmark
@Group("Hierarchy")
public double testHierarchy(HierarchyState state) {
    long sum = 0;
    for (int i = 0; i < state.array.length; i++) {
        sum += (state.array[i]).x + (state.array[i]).y; // hierarchy.x + hierarchy.y
    }
    return sum;
}
```

```
/** array.length = 2048 */
```

Fields Layout

Benchmark comparison

	JIT C2 us/op	LLVM Clang us/op
Hierarchy	1.96	1.93

lower is better

Fields Layout

Conclusions

Java automatically does fields reordering, nevertheless the object header impacts dCache.

LLVM clang does not automatically do fields reordering (i.e. it is at the mercy of developers).

**Cases where fields reordering helps and
where we should pay attention!**

Fields Reordering

Useful

- **improves the overall object size** by minimizing the object layout
- **prevents unaligned fields access** (i.e. objects - 8 bytes alignment / fields – type aligned)
- **re-grouping fields** so they might fall on the same cache line, hence **minimizing CPU cache misses** on fields access

Fields Reordering

Useful

- **improves the overall object size** by minimizing the object layout
- **prevents unaligned fields access** (i.e. objects - 8 bytes alignment / fields - type aligned)
- **re-grouping fields** so they might fall on the same cache line, hence **minimizing CPU cache misses** on fields access

Pay Attention

- **false-sharing** problem
- **re-grouping fields** might also be problematic when they end up in falling on different cache lines, **increasing CPU cache misses**

Null Checks

Null Checks

Source code

```
@Benchmark
public int hotMethod(Integer anInt) { // anInt != NULL
    int counter = 0;

    if (null != anInt)
        counter = anInt * 42;

    return counter;
}
```

Null Checks

JIT C2

```
mov    r10d,DWORD PTR [rsi+0xc]      ; *getfield anInt
imul   eax,DWORD PTR [r12+r10*8+0xc],0x2a ; 0x2a = 42
; implicit exception: dispatches to 0x00007f02e038c062
```

implicit NULL check handler



Null Checks

JIT C2

```
mov     r10d,DWORD PTR [rsi+0xc]           ; *getfield anInt
imul    eax,DWORD PTR [r12+r10*8+0xc],0x2a ; 0x2a = 42
                                              ; implicit exception: dispatches to 0x00007f02e038c062
```

implicit NULL check handler

```
0x00007f02e038c062: mov     esi,0xffffffff5d
                  mov     DWORD PTR [rsp],r10d
                  call    0x00007f02e0317c00 ; {runtime_call UncommonTrapBlob}
                  call    0x00007f02e4bae300 ; *if_acmpeq {reexecute=0 rethrow=0 return_oop=0}
                                              ; - com.jpt.NullSanityChecks::nullCheck@4
```

uncommon trap

Null Checks

LLVM clang

explicit NULL check



```
mov rax, qword ptr [rip + x]
test rax, rax
je L0
imul eax, dword ptr [rax], 42
ret
```

L0

```
xor eax, eax
ret
```


Null Checks

Benchmark comparison

JIT C2 ns/op	LLVM Clang ns/op
1.8	2

lower is better

Explicit NULL check in this particular test case does not make a big difference between these two approaches on modern hardware.

[note] could not find any way for C++ Google benchmark to have ns decimal precision

Null Checks

Conclusions

LLVM clang adds an explicit NULL check (i.e. `hotMethod()` has the explicit check).

JIT C2 relies on signals (e.g. `SEG_FAULT`) natively provided by hardware and handled by OS to manage NULLs, it optimizes without adding the NULL check .

**What happens if hotMethod() is called
with both NULL and !NULL ?**

Null Checks

Source code

```
@Benchmark
public int hotMethod(Integer anInt) { // anInt == NULL || anInt != NULL
    int counter = 0;


    if (null != anInt)
        counter = anInt * 42;

    return counter;
}
```

Null Checks

JIT C2

explicit NULL check




```
mov     r11d, DWORD PTR [rsi+0xc]           ; *getfield anInt
test    r11d, r11d
je      L0
imul    eax, DWORD PTR [r12+r11*8+0xc], 0x2a ; 0x2a = 42
jmp     L1
L0      xor     eax, eax
L1      add     rsp, 0x10
```

Null Checks

JIT C2

explicit NULL check



```
mov     r11d, DWORD PTR [rsi+0xc]           ; *getfield anInt
test    r11d, r11d
je      L0
imul    eax, DWORD PTR [r12+r11*8+0xc], 0x2a ; 0x2a = 42
jmp     L1

L0      xor     eax, eax
L1      add     rsp, 0x10
```

Not any **uncommon trap**, JIT falls back to a classic way of handling NULLs (similar to LLVM clang)

Null Checks

Benchmark comparison

JIT C2 ns/op	LLVM Clang ns/op
2.2	2
<i>lower is better</i>	

“Fast enough is fast enough” - Cliff Click

Null Checks

When the **number of well-predictable branches** within a code region **goes beyond the hardware limit** (i.e. total amount of branch predictor resource is finite), **branch prediction starts falling off the cliff and there might be a significant difference between these two!**

Lock Elision

Lock Elision

Source code

```
public int hotMethod() {  
    int sum = 0;  
  
    Object lock = new Object();  
  
    synchronized (lock) {  
        sum += value;  
    }  
  
    // ...  
  
    synchronized (lock) {  
        sum += value;  
    }  
  
    return sum;  
}
```



10 synchronized blocks

Lock Elision

JIT C2

```
mov    r11d,DWORD PTR [rsi+0xc]  ;*getfield value
mov    eax,r11d
add    eax,r11d
add    eax,r11d
add    eax,r11d
add    eax,r11d
add    eax,r11d
add    eax,r11d
add    eax,r11d
add    eax,r11d
add    eax,r11d
add    eax,r11d
```

Lock Elision

JIT C2

sequential sum
- no locks in between -

```
mov    r11d,DWORD PTR [rsi+0xc]    ;*getfield value
mov    eax,r11d
add    eax,r11d
add    eax,r11d
add    eax,r11d
add    eax,r11d
add    eax,r11d
add    eax,r11d
add    eax,r11d
add    eax,r11d
add    eax,r11d
```

Lock Elision

LLVM clang

```
mov r14d, edi
```

```
call std::__1::mutex::lock()
```

```
mov rdi, rbx
```

```
call std::__1::mutex::unlock()
```

```
// ...
```

```
mov rdi, rbx
```

```
call std::__1::mutex::lock()
```

```
add r14d, r14d
```

```
lea ebp, [r14 + 4*r14]
```

```
mov rdi, rbx
```

```
call std::__1::mutex::unlock()
```

9 useless locks

```
; r14d <- r14d + r14d
```

```
; ebp <- [r14 + 4 x r14]
```

Lock Elision

LLVM clang

```
mov r14d, edi
```

```
call std::__1::mutex::lock()  
mov rdi, rbx  
call std::__1::mutex::unlock()
```

```
// ...
```

```
mov rdi, rbx  
call std::__1::mutex::lock()  
add r14d, r14d  
lea ebp, [r14 + 4*r14]  
mov rdi, rbx  
call std::__1::mutex::unlock()
```


9 useless locks

compute sum inside last
(10th) synchronized block

```
; r14d <- r14d + r14d  
; ebp <- [r14 + 4 x r14]
```

Lock Elision

Benchmark comparison



Lock Elision

Conclusions

LLVM clang cannot optimize the locks, however the additions are done inside last synchronized block.

Runtime performance is dependent on number of locks.

Java's built-in synchronized semantics are well-defined by the language hence JIT C2 takes advantage of them and remove locks which does not escape.

Runtime performance is not dependent on number of locks.

Virtual Calls

Single Target Invocation


Source code

```
@Param({"3"})
public int param;

@Setup
public void setup() {
    shape = new Triangle();
}

public int compute(Shape shape, int i) {
    return shape.compute(i); // return (param * i)
}

@Benchmark
@Group("Monomorphic")
public int testCompute(State state) {
    return compute(state.shape, param); // return (3 * 17)
}
```



Single Target Invocation

JIT C2

```
mov    eax,ecx  
shl    eax,0x4    ; eax <- eax x 16  
add    eax,ecx    ; eax <- param x 17
```

Single target invocation (e.g. monomorphic call) is inlined, no any **uncommon trap**.

Single Target Invocation

LLVM clang

```
mov rdi, qword ptr [rip + (anonymous namespace)::triangle]
mov rax, qword ptr [rip + (anonymous namespace)::param]
mov esi, dword ptr [rax]
mov rax, qword ptr [rdi]
mov rax, qword ptr [rax + 16] ; rax <- entry from vtable e.g. compute()
jmp rax
```



Passes the control to the
virtual call

Single Target Invocation

LLVM clang

```
mov rdi, qword ptr [rip + (anonymous namespace)::triangle]  
mov rax, qword ptr [rip + (anonymous namespace)::param]  
mov esi, dword ptr [rax]  
mov rax, qword ptr [rdi]  
mov rax, qword ptr [rax + 16] ; rax <- entry from vtable e.g. compute()  
jmp rax
```



Passes the control to the
virtual call

Single target invocation uses vtable to perform the virtual call.

Single Target Invocation

Benchmark comparison

	JIT C2 ns/op	LLVM Clang ns/op
Monomorphic calls	1.9	2

lower is better


[note] could not find any way for C++ Google benchmark to have ns decimal precision

Two Target Invocation

Source code

```
@Setup
public void setup() {
    shape1 = new Triangle();
    shape2 = new Rectangle();
}

@Benchmark
@Group("Bimorphic")
public int testCompute(State state) {
    return compute(state.shape1, param) + compute(state.shape2, param);
    // return (3 * 17) + (3 * 19)
}
```



Two Target Invocation

JIT C2

```
mov     r10d,DWORD PTR [rdx+0x8] ; implicit exception: dispatches to 0x00007fe7cdc03991
cmp     r10d,0xf801f06b          ; {metadata(com/jpt/Rectangle)}
je      L0
cmp     r10d,0xf801f02c          ; {metadata(com/jpt/Triangle)}
jne     0x00007fe7cdc0397a       ; *invokevirtual compute (e.g. !Triangle && !Rectangle)
mov     eax,ecx
shl     eax,0x4
add     eax,ecx                  ; e.g. eax <- param x 17
jmp     L1
L0:     imul    eax,ecx,0x13       ; e.g. eax <- param x 19
L1:     add     rsp,0x20
pop     rbp
```

0x00007fe7cdc0397a: call 0x00007fe7cdb8ac00 ; *invokevirtual compute()



uncommon trap

Two Target Invocation

JIT C2

```
mov     r10d,DWORD PTR [rdx+0x8] ; implicit exception: dispatches to 0x00007fe7cdc03991
cmp     r10d,0xf801f06b          ; {metadata(com/jpt/Rectangle)}
je      L0
cmp     r10d,0xf801f02c          ; {metadata(com/jpt/Triangle)}
jne     0x00007fe7cdc0397a       ; *invokevirtual compute (e.g. !Triangle && !Rectangle)
mov     eax,ecx
shl     eax,0x4
add     eax,ecx                  ; e.g. eax <- param x 17
jmp     L1
L0:     imul    eax,ecx,0x13       ; e.g. eax <- param x 19
L1:     add     rsp,0x20
pop     rbp
```

0x00007fe7cdc0397a: call 0x00007fe7cdb8ac00 ; *invokevirtual compute()



uncommon trap

Two targets call (e.g. bimorphic call) is also inlined but **uncommon trap** is added.

Two Target Invocation

LLVM clang

```
mov rdi, qword ptr [rip + (anonymous namespace)::triangle]
mov rax, qword ptr [rip + (anonymous namespace)::param]
mov esi, dword ptr [rax]
mov rax, qword ptr [rdi]
call qword ptr [rax + 16] ; e.g. compute() virtual call
mov ebx, eax              ; ebx <- param x 17

mov rdi, qword ptr [rip + (anonymous namespace)::rectangle]
mov rax, qword ptr [rip + (anonymous namespace)::param]
mov esi, dword ptr [rax]
mov rax, qword ptr [rdi]
call qword ptr [rax + 16] ; e.g. compute() virtual call

add eax, ebx              ; eax <- (param x 17) + (param x 19)
```

Two Target Invocation

LLVM clang

Triangle::compute(int)

```
mov rdi, qword ptr [rip + (anonymous namespace)::triangle]  
mov rax, qword ptr [rip + (anonymous namespace)::param]  
mov esi, dword ptr [rax]  
mov rax, qword ptr [rdi]  
call qword ptr [rax + 16] ; e.g. compute() virtual call  
mov ebx, eax ; ebx <- param x 17
```

```
mov rdi, qword ptr [rip + (anonymous namespace)::rectangle]  
mov rax, qword ptr [rip + (anonymous namespace)::param]  
mov esi, dword ptr [rax]  
mov rax, qword ptr [rdi]  
call qword ptr [rax + 16] ; e.g. compute() virtual call
```

```
add eax, ebx ; eax <- (param x 17) + (param x 19)
```

Two Target Invocation

LLVM clang

Triangle::compute(int)

```
mov rdi, qword ptr [rip + (anonymous namespace)::triangle]  
mov rax, qword ptr [rip + (anonymous namespace)::param]  
mov esi, dword ptr [rax]  
mov rax, qword ptr [rdi]  
call qword ptr [rax + 16] ; e.g. compute() virtual call  
mov ebx, eax ; ebx <- param x 17
```

Rectangle::compute(int)

```
mov rdi, qword ptr [rip + (anonymous namespace)::rectangle]  
mov rax, qword ptr [rip + (anonymous namespace)::param]  
mov esi, dword ptr [rax]  
mov rax, qword ptr [rdi]  
call qword ptr [rax + 16] ; e.g. compute() virtual call  
  
add eax, ebx ; eax <- (param x 17) + (param x 19)
```

Two Target Invocation

LLVM clang

Triangle::compute(int)

```
mov rdi, qword ptr [rip + (anonymous namespace)::triangle]  
mov rax, qword ptr [rip + (anonymous namespace)::param]  
mov esi, dword ptr [rax]  
mov rax, qword ptr [rdi]  
call qword ptr [rax + 16] ; e.g. compute() virtual call  
mov ebx, eax ; ebx <- param x 17
```

Rectangle::compute(int)

```
mov rdi, qword ptr [rip + (anonymous namespace)::rectangle]  
mov rax, qword ptr [rip + (anonymous namespace)::param]  
mov esi, dword ptr [rax]  
mov rax, qword ptr [rdi]  
call qword ptr [rax + 16] ; e.g. compute() virtual call
```

```
add eax, ebx ; eax <- (param x 17) + (param x 19)
```

Two target invocation uses vtable to perform the virtual calls.

Two Target Invocation

Benchmark comparison

	JIT C2 ns/op	LLVM Clang ns/op
Bimorphic calls	3.9	4

lower is better


[note] could not find any way for C++ Google benchmark to have ns decimal precision

Three Target Invocation

Source code

```
@Setup
public void setup() {
    shape1 = new Triangle();
    shape2 = new Rectangle();
    shape3 = new Square();
}

@Benchmark
@Group("Megamorphic")
public int testCompute(State state) {
    return compute(state.shape1, param) + compute(state.shape2, param) + compute(state.shape3, param);
    // return (3 * 17) + (3 * 19) + (3 * 23)
}
```



Three Target Invocation

JIT C2

```
mov    rsi,rdx
mov    edx,ecx
mov    rax,0xffffffffffffffff
call   0x00007fe469b89500 ; *invokevirtual compute()
```



virtual call

Megamorphic call is not anymore inlined, neither any **uncommon trap** is added.

Three Target Invocation

LLVM clang

```
mov rdi, qword ptr [rip + (anonymous namespace)::triangle]
mov rax, qword ptr [rip + (anonymous namespace)::param]
mov esi, dword ptr [rax]
mov rax, qword ptr [rdi]
call qword ptr [rax + 16] ; e.g. compute() virtual call
mov ebx, eax              ; ebx <- param x 17
```

```
mov rdi, qword ptr [rip + (anonymous namespace)::rectangle]
mov rax, qword ptr [rip + (anonymous namespace)::param]
mov esi, dword ptr [rax]
mov rax, qword ptr [rdi]
call qword ptr [rax + 16] ; e.g. compute() virtual call
mov ebp, eax              ; ebp <- param x 19
```

```
add ebp, ebx              ; ebp <- (param x 17) + (param x 19)
```

```
mov rdi, qword ptr [rip + (anonymous namespace)::square]
mov rax, qword ptr [rip + (anonymous namespace)::param]
mov esi, dword ptr [rax]
mov rax, qword ptr [rdi]
call qword ptr [rax + 16] ; e.g. compute() virtual call
```

```
add eax, ebp ; ebp <- (param x 17) + (param x 19) + (param x 23)
```

Three Target Invocation

LLVM clang

Triangle::compute(int)

```
mov rdi, qword ptr [rip + (anonymous namespace)::triangle]
mov rax, qword ptr [rip + (anonymous namespace)::param]
mov esi, dword ptr [rax]
mov rax, qword ptr [rdi]
call qword ptr [rax + 16] ; e.g. compute() virtual call
mov ebx, eax              ; ebx <- param x 17
```

```
mov rdi, qword ptr [rip + (anonymous namespace)::rectangle]
mov rax, qword ptr [rip + (anonymous namespace)::param]
mov esi, dword ptr [rax]
mov rax, qword ptr [rdi]
call qword ptr [rax + 16] ; e.g. compute() virtual call
mov ebp, eax              ; ebp <- param x 19
```

```
add ebp, ebx              ; ebp <- (param x 17) + (param x 19)
```

```
mov rdi, qword ptr [rip + (anonymous namespace)::square]
mov rax, qword ptr [rip + (anonymous namespace)::param]
mov esi, dword ptr [rax]
mov rax, qword ptr [rdi]
call qword ptr [rax + 16] ; e.g. compute() virtual call
```

```
add eax, ebp ; ebp <- (param x 17) + (param x 19) + (param x 23)
```

Three Target Invocation

LLVM clang

Triangle::compute(int)

```
mov rdi, qword ptr [rip + (anonymous namespace)::triangle]
mov rax, qword ptr [rip + (anonymous namespace)::param]
mov esi, dword ptr [rax]
mov rax, qword ptr [rdi]
call qword ptr [rax + 16] ; e.g. compute() virtual call
mov ebx, eax                ; ebx <- param x 17
```

Rectangle::compute(int)

```
mov rdi, qword ptr [rip + (anonymous namespace)::rectangle]
mov rax, qword ptr [rip + (anonymous namespace)::param]
mov esi, dword ptr [rax]
mov rax, qword ptr [rdi]
call qword ptr [rax + 16] ; e.g. compute() virtual call
mov ebp, eax                ; ebp <- param x 19
```

```
add ebp, ebx                ; ebp <- (param x 17) + (param x 19)
```

```
mov rdi, qword ptr [rip + (anonymous namespace)::square]
mov rax, qword ptr [rip + (anonymous namespace)::param]
mov esi, dword ptr [rax]
mov rax, qword ptr [rdi]
call qword ptr [rax + 16] ; e.g. compute() virtual call
```

```
add eax, ebp ; ebp <- (param x 17) + (param x 19) + (param x 23)
```

Three Target Invocation

LLVM clang

Triangle::compute(int)

```
mov rdi, qword ptr [rip + (anonymous namespace)::triangle]
mov rax, qword ptr [rip + (anonymous namespace)::param]
mov esi, dword ptr [rax]
mov rax, qword ptr [rdi]
call qword ptr [rax + 16] ; e.g. compute() virtual call
mov ebx, eax              ; ebx <- param x 17
```

Rectangle::compute(int)

```
mov rdi, qword ptr [rip + (anonymous namespace)::rectangle]
mov rax, qword ptr [rip + (anonymous namespace)::param]
mov esi, dword ptr [rax]
mov rax, qword ptr [rdi]
call qword ptr [rax + 16] ; e.g. compute() virtual call
mov ebp, eax              ; ebp <- param x 19
```

```
add ebp, ebx              ; ebp <- (param x 17) + (param x 19)
```

Square::compute(int)

```
mov rdi, qword ptr [rip + (anonymous namespace)::square]
mov rax, qword ptr [rip + (anonymous namespace)::param]
mov esi, dword ptr [rax]
mov rax, qword ptr [rdi]
call qword ptr [rax + 16] ; e.g. compute() virtual call
```

```
add eax, ebp ; ebp <- (param x 17) + (param x 19) + (param x 23)
```

Three Target Invocation

LLVM clang

Triangle::compute(int)

```
mov rdi, qword ptr [rip + (anonymous namespace)::triangle]
mov rax, qword ptr [rip + (anonymous namespace)::param]
mov esi, dword ptr [rax]
mov rax, qword ptr [rdi]
call qword ptr [rax + 16] ; e.g. compute() virtual call
mov ebx, eax              ; ebx <- param x 17
```

Rectangle::compute(int)

```
mov rdi, qword ptr [rip + (anonymous namespace)::rectangle]
mov rax, qword ptr [rip + (anonymous namespace)::param]
mov esi, dword ptr [rax]
mov rax, qword ptr [rdi]
call qword ptr [rax + 16] ; e.g. compute() virtual call
mov ebp, eax              ; ebp <- param x 19
```

```
add ebp, ebx              ; ebp <- (param x 17) + (param x 19)
```

Square::compute(int)

```
mov rdi, qword ptr [rip + (anonymous namespace)::square]
mov rax, qword ptr [rip + (anonymous namespace)::param]
mov esi, dword ptr [rax]
mov rax, qword ptr [rdi]
call qword ptr [rax + 16] ; e.g. compute() virtual call
```

```
add eax, ebp ; ebp <- (param x 17) + (param x 19) + (param x 23)
```

Three Target Invocation

LLVM clang

Triangle::compute(int)

```
mov rdi, qword ptr [rip + (anonymous namespace)::triangle]
mov rax, qword ptr [rip + (anonymous namespace)::param]
mov esi, dword ptr [rax]
mov rax, qword ptr [rdi]
call qword ptr [rax + 16] ; e.g. compute() virtual call
mov ebx, eax              ; ebx <- param x 17
```

Rectangle::compute(int)

```
mov rdi, qword ptr [rip + (anonymous namespace)::rectangle]
mov rax, qword ptr [rip + (anonymous namespace)::param]
mov esi, dword ptr [rax]
mov rax, qword ptr [rdi]
call qword ptr [rax + 16] ; e.g. compute() virtual call
mov ebp, eax              ; ebp <- param x 19
```

```
add ebp, ebx              ; ebp <- (param x 17) + (param x 19)
```

Square::compute(int)

```
mov rdi, qword ptr [rip + (anonymous namespace)::square]
mov rax, qword ptr [rip + (anonymous namespace)::param]
mov esi, dword ptr [rax]
```

Three target invocation uses vtable to perform the virtual calls.

```
add eax, ebp ; ebp <- (param x 17) + (param x 19) + (param x 23)
```

Three Target Invocation

Benchmark comparison

	JIT C2 ns/op	LLVM Clang ns/op
Megamorphic calls	8.1	7

lower is better

[note] could not find any way for C++ Google benchmark to have ns decimal precision

Three Target Invocation

Conclusions

LLVM clang does use `vtable` to perform virtual calls (**in this case!**).

Runtime performance is dependent on number of target implementations.

JIT C2 is able to inline (e.g. monomorphic and bimorphic calls) and to add uncommon traps (e.g. bimorphic calls).

Starting with 3th implementation it just performs a virtual call → due to “`polluted profile`” context (i.e. method is used in many different contexts with independent operand types).

Profile pollution

Profiles (especially [type profiles](#)) are subject to pollution if the profiled code is heavily reused in ways that diverge from each other.

As a simple example, if `ArrayList.contains` is used with lists that never contain nulls, some null checks will never be taken, and the profile can reflect this. But if this routine is also used with lists that occasionally contain nulls, then the "taken" count of the null check instruction may become non-zero. This in turn may influence the compiler to check operands more cautiously, with a loss of performance for *all* uses of the method.

As a more complex example, if `ArrayList.contains` is only ever used on arrays that contain strings, then the type profile will reflect this, and the virtual call to `Object.equals` on each element can be optimistically inlined as if it were an `invokespecial` of `String.equals`.

Polluted profiles can be mitigated by a number of means, including:

- inlining with type or value information flowing from caller context
- inlining more than one type case (`UseBimorphicInlining`)
- [context-dependent split profiles](#) (bug 8015416)
- hand inlining by the Java programmer (which is discouraged)
- generic type reification (when pollution comes from a type parameter; not implemented)
- other forms of on-line code splitting
- adding type-check bytecodes for profiled references *before* the call to the shared routine (including `instanceof Object` which is a no-op except for profile effects)



JDK / JDK-8015416

tier one should collect context-dependent split profiles

Details

Type:	Enhancement	Status:	OPEN
Priority:	P3	Resolution:	Unresolved
Affects Version/s:	9, 10	Fix Version/s:	13
Component/s:	hotspot		
Labels:	c1 comp-backlog performance tiered-compilation		
Subcomponent:	compiler		

Description

To avoid polluted profiles, tier one should create disjoint "split" profiles for methods that it inlines. If a method is inlined three times, it should get three fresh copies of its global profile, initialized to no types or counts. When a later tier re-optimizes the code and inlines the same method, the compiler should use the context-dependent profile in preference to the global profile, if one is available.

Scalar Replacement

Scalar Replacement

Source code

```
@Param({"3"})
private int param1;

@Param({"5"})
private int param2;

@Benchmark
public int hotMethod() {
    Wrapper w = new Wrapper(param1, param2);
    return w.x + w.y;
}

public static class Wrapper {
    public int x;
    public int y;

    public Wrapper(int value1, int value2) {
        this.x = value1;
        this.y = value2;
    }
}
```

Scalar Replacement

JIT C2

```
mov    eax,DWORD PTR [rsi+0x10]  
add    eax,DWORD PTR [rsi+0xc]    ; eax <- x + y
```

Scalar Replacement

LLVM clang

```
mov rax, qword ptr [rip + (anonymous namespace)::x]  
mov rcx, qword ptr [rip + (anonymous namespace)::y]  
mov esi, dword ptr [rcx]  
add esi, dword ptr [rax]      ; eax <- x + y
```

Scalar Replacement

Benchmark comparison

JIT C2 ns/op	LLVM Clang ns/op
1.8	2

lower is better

“Fast enough is fast enough” - Cliff Click

Scalar Replacement

Conclusions

LLVM clang does constant folding at compile time, avoiding allocation.

JIT C2 prevents the heap memory allocations (i.e. NoEscape object state), as well.

Conclusions

LLVM clang can trigger - even ahead of time - a lot of powerful optimizations (e.g. *sum 1 to N integers, scalar replacement*).

JIT C2 can do better optimizations than LLVM clang when runtime profiling information matters or built-in synchronized semantics are well-defined at language level (e.g. *lock elision, virtual calls*).


When both compilers optimize the code based on similar approaches (e.g. *loop unrolling/vectorization, null checks, sum(array[], array[]), sum(array[], const)*) there is almost identical performance.

In Closing ...

“It doesn't make a difference how beautiful your guess is [...] how smart you are, who made the guess, or what his name is.

If it disagrees with experiment, it's wrong.”

The Richard Feynman Lectures on Physics: The New Millennium Edition



Special thanks to

George Frăsineanu

C++ Senior Developer
george.frasineanu@gmail.com



THANK YOU

Ionuț Baloșin

Software Architect

www.ionutbalosin.com



[@ionutbalosin](https://twitter.com/ionutbalosin)



Krakow, 20-22 June 2018

@ionutbalosin

Full Tests Configuration

Hardware

- Intel i7-6700HQ Skylake
- 16GB DDR4 2133 MHz



Software

- Ubuntu 16.04.2
- JDK 10
- Java Measurement Harness
- JITWatch
- LLVM clang 6 (w/ -O3 optimization level)
- C++ Google Benchmark
- godbolt.org
- perf





THANK YOU

Ionuț Baloșin

Software Architect

www.ionutbalosin.com



[@ionutbalosin](https://twitter.com/ionutbalosin)



Krakow, 20-22 June 2018

@ionutbalosin